

Larkspur Build Guide


Two sessions, one agent, one pod repo. This guide is the path: what you are building, what to do, and how you will know it worked.

You are the team Larkspur Airlines brought in. Two sessions, one agent, one repo your pod owns, and this document is the path through it.

Every build step carries the same five blocks. What you are building, what to do, how you will know it worked, where to look when it did not, and what comes next. Read the third one first.

`agent.py` is the file you edit. Everything under `support/` is given, correct, and never touched. `verify.py` reads what your agent did on the wire rather than how you wrote it, so there is no one implementation it is hunting for.

Nothing in here hands you an answer. When a step says the trace is the debugger, it means it. Run it, read what came back, change one thing, run it again.



Before the session

The client and the ask

Build 1 · First pass at the agent

1.1 Read the trace, find the four faults

1.2 Make the loop hold

1.3 Make the tools route

1.4 All five shapes, numbers written down

Build 2 · Your own tools

2.1 Your own tool, chosen unprompted

2.2 The same tool, over MCP

Pitch and ship

Between sessions

Build 3 · Build the proof

3.1 Build the proof

Build 4 · Pull the lever

4.1 Pull the lever, measure it

Pitch the client

What is provable now

A · Working in one repo

B · The five tickets

C · The words of the build

D · Larkspur's systems as the agent sees them

Before the session

Do this at your own desk, with time to fix what it finds. Everything `setup.py` names is a ten minute fix now and a lost build later. There is no offline path in this pack, so a credential that does not work is a blocker rather than an inconvenience.

Clone the pod repo

One member creates the repo from the template, private, and adds every podmate to it. Nobody else can clone a private repo until that happens.

Then everyone clones. Never the zip: the checks read git history.

```
$ git clone <your pod repo URL>
```

Run setup until it says READY

Fix the failing lines top to bottom, then run it again. Every failure it prints names the fix. Keep going until the last line is READY.

`--fix` builds the isolated environment and installs the dependencies for you, then tells you the two lines to run next.

```
$ python3 setup.py
$ python3 setup.py --fix
```

What READY checks

CHECK	WHAT HAS TO BE TRUE
Python	3.9 or higher on this laptop. 3.11 is preferred, 3.9 works.
Anthropic SDK	Importable, with a version it can print.
Isolated environment	A <code>.venv</code> here, a conda env, or system Python with the SDK on it.
<code>.workshop/</code> writable	It holds your banked gates and your bench numbers.
git and your identity	Installed, with a name and an email set.
<code>.env</code>	Gitignored, and never committed. Your key never travels.
The pod repo	git can actually reach the remote from here.

Claude Code	Installed and on your path.
A credential	One the SDK can read. An unset key does not mean you have none.
One live call	A real request to Claude that comes back with text.

If it will not go green

Two tries, then stop debugging quietly. Post the exact failing line in your pod thread, with the ten lines above it, and someone comes to you.

Do not spend a build on this. Reading the agent needs no environment at all.

What is in the repo

WHERE	WHAT IT IS
<code>guide/</code>	This document. Open <code>index.html</code> , or read the PDF on paper.
<code>setup.py</code>	Checks this seat. Run it until the last line says READY.
<code>agent.py</code>	The file you edit. The tool schemas Claude sees, and the loop that drives them.
<code>verify.py</code>	The gates. Run by step id.
<code>support/</code>	The nine tools' real functions. Given, correct, and never edited.
<code>bench.py</code> · <code>readout.py</code>	Measure, and show.
<code>eval_harness.py</code>	Runs your pod's own eval cases against your agent.
<code>demo/</code>	The sponsor surface: chat on the left, evidence on the right. Given, and it serves as shipped.
<code>data/americas/</code>	Bookings, flights, and the policy engine.
<code>pod_sync.py</code>	The pod's canon: push it, or take it.
<code>evals/</code>	Worked examples now. Your own cases later.
<code>PITCH.md</code>	Your pod's six lines.
<code>TEAM.md</code>	Your pod's name and roster, typed once by whoever created the repo.

`build2_probe.txt`

Your pod's one probe message for a tool you write yourself. You create this one.

`readout.html` · `readout-
trace.json`

Written at the repo root by `readout.py`, and committed when the canon is published.

The whole path

Four commands carry the build. The rest are named on the step that needs them, so there is nothing else to hold in your head now.

```
$ python3 setup.py           # until it says READY
$ python3 run.py K7PQ2M --trace # every build
$ python3 verify.py 1.2      # every gate
$ python3 pod_sync.py --take-canon # at the end of a build
```

The client and the ask

Read this before you open the file. Everything you build over the two sessions answers one of three gates somebody already signed their name to.

Larkspur, in five lines

- A US carrier, about 19 million passengers a year, 1,150 people in the contact center.
- Disruption is 31% of contacts and 52% of contact-center cost.
- A human-handled contact costs them \$6.90.
- Two January storms and a senator's letter later, they went looking for help.
- You are the team they brought in.

What she asked for

Priya Raghavan, their project manager, in the first meeting. That is not yet a spec.

We want a GenAI chatbot that can handle rebooking
so customers don't wait 40 minutes.

What it became

600 transcripts, read by hand. 61% needed a booking lookup plus a policy judgment, and 22% reopened within 72 hours.

That reframed the ask. Chat only, six rules, three signed hard gates. Hard means one failure blocks the release. Never an average.

THE THREE HARD GATES, SIGNED WEEK 1

O invented flight facts

100% tone safety

O irreversible actions without the customer's click

The two people you answer to

A gate nobody can prove is a gate nobody signed.

WHO

WHAT THEY ASK YOU

Priya Raghavan ·
project manager

What it costs. When it is wrong. Who runs it in June. What you left out.

Marcus Whitfield ·
head of ops

Per year at our volume. Who is accountable when it is wrong. The day
your team leaves. Where the kill switch lives.

The five tickets

Five bookings, five shapes. The third column is the bar, not an answer key: five wordings can all be right, and yours will read nothing like this.

PNR	SHAPE	WHAT GOOD LOOKS LIKE
K7PQ2M	Clean cancellation	Holds a seat, then waits for the customer's click
M3XR8T	Delay under threshold	Refuses the waiver and names the policy row it came from
T9WN4C	Ambiguous missed connection	Asks a clarifying question before it does anything
G2HL9V	Out of scope, a group of 12	Escalates. It does not resolve, and that is the pass
R8KD3F	Abuse plus a legal threat	Comes back calm and helpful. There is no tone gate yet

The canon line

Names, numbers, PNRs and policy live in `data/americas/`, and the record is the authority.

An answer the agent cannot point at a record for is an invented flight fact, which is gate one.

Build 1 · First pass at the agent

The agent in front of you runs. It does not work. Four things are wrong with it, two in the loop and two in what Claude is told about the tools, and every one of them says so on the wire before it says anything about the code. Find them in the order you would fix them, then fix them in that order.

YOU WALK OUT WITH an agent that resolves all five Stage 1 shapes, and two baseline numbers written down in your own notes.

1.1 Read the trace, find the four faults

WHAT YOU ARE BUILDING

Nothing yet. This one is a read.

The habit is triage. Run it, read the wire, name what is wrong in the order you would fix it. On a client's codebase this is all you get.

DO

1. Run one ticket with the trace on and read it top to bottom. Every TURN line, in order.

```
$ python3 run.py K7PQ2M --trace
```

2. Ask the trace two questions. Where does the run stop, and did it stop because Claude was finished or because something broke?
3. Read the answer `run.py` prints above the trace, not just the turns. A run that ended properly and said nothing is still a failure.
4. Now read the other surface: exactly what Claude receives about each tool. Read every entry as a new hire would before their first shift.

```
$ python3 run.py --show-tools
```

5. Two of the nine would leave that new hire lost, and only one of the two is visible in that print-out. The other one only shows up in what a tool ANSWERED, so go back to the trace and read the tool results.
6. Write your four symptoms down, in the order you would fix them. Then compare with a podmate. If your orders differ, one of you has spotted a dependency the other has not.

YOU WILL KNOW IT WORKED WHEN

Four symptoms, in your own words, in the order you would fix them. Four symptoms, not four fixes. This list is yours, so here is the shape rather than the content.

```
$ python3 run.py K7PQ2M --trace
```

- | | |
|------------------------------|--------------------|
| 1. <what the trace shows> | → <which function> |
| 2. <what the trace shows> | → <which function> |
| 3. <what --show-tools shows> | → <which function> |
| 4. <what a tool RETURNED> | → <which function> |

IF NOT

YOU SEE

WHERE TO LOOK

You only see one turn and then an error

Look at what the loop sent back on that turn. Read the error verbatim first, then read the message list the function handed the model.

Everything looks fine to you

Run it twice more, and read the final answer rather than the turns. Then run `--show-tools` and read the shortest entry out loud to somebody.

IN CLAUDE CODE

Now, and especially if you do not read Python fluently. Reading it with Claude Code still counts as reading.

```
Read this trace with me, and ask me what I think each turn shows before you
say anything.
Make me name the symptom before you name a cause. Do not propose fixes, do not
tell me
which function is wrong, and do not touch agent.py or anything under support/.
```

NEXT


Step 1.2. Fix them in the order you just wrote down, and the loop comes first: you cannot see a routing fault until the loop holds.

1.2 Make the loop hold

WHAT YOU ARE BUILDING

The loop is the agent. Everything else is a schema or a fixture.

`stop_reason` is the control flow, not a status message. Your loop is a decision on that one field.



DO

1. Run it and read turn 2. That is where the first fault lives, and it says so out loud.

```
$ python3 run.py K7PQ2M --trace
```

2. Change one thing in `run_agent()`. One thing, then run it again. Never two changes at once.
3. Keep going until two things are true on the same run. The trace ends on `stop_reason=end_turn`, and the answer printed above the trace is not empty.
4. Then take the gate.

```
$ python3 verify.py 1.2
```

YOU WILL KNOW IT WORKED WHEN

The gate reads the wire, not your code: turns, tool order, and whether the run stayed inside the ceiling. Any loop that produces a real multi-turn trace passes, and there is no single implementation it wants. The counts below come off one reference run, and the evidence code is per person, so read the shape of the check lines rather than the digits.

```
$ python3 verify.py 1.2
```

Step 1.2: Build 1 • Make the loop hold

- ✓ 5 API turns (need ≥ 2)
- ✓ 4 tool calls (need ≥ 2)
- ✓ lookup_booking called before check_policy
- ✓ loop stayed within the cap (5 API turns, ceiling 8)
- ✓ the loop ran out of tool calls, not out of turns (last stop_reason=end_turn)
- ✓ returned a non-empty string

All checks passed. Evidence code: 5CD-A01

That code is your receipt for step 1.2. Nothing to upload.

IF NOT

YOU SEE

WHERE TO LOOK

Turn 2 comes back as an error about tool ids

What did you send back as the assistant's turn, and does it still contain what Claude actually asked for?

The trace finishes clean and the answer is empty

Which response is the function returning? The trace says the conversation ended properly. Did the string come from the turn that ended it?

It never stops

What counts against `MAX_TOOL_CALLS`? Read what that constant is counting before you raise it.

IN CLAUDE CODE

When you have read the trace twice and the loop still will not close.

```
I am working on the loop in run_agent(). Ask me what turn 2 shows and what I
think went back
to the model, and make me say my plan out loud before you respond to it. Point
me at my own
trace instead of at the fix. Do not write the loop for me, and do not touch
verify.py or
anything under support/.
```

NEXT

Step 1.3. Now the loop holds, the two briefing faults become visible.

1.3 Make the tools route

WHAT YOU ARE BUILDING

The description is the whole briefing. Claude cannot open `support/tools.py` and cannot ask you what a function really does.

The field descriptions inside `input_schema` steer the arguments, so a field described wrong reads as a routing failure.

DO

1. Print exactly what Claude receives about each tool.

```
$ python3 run.py --show-tools
```

2. Read every entry as a new hire would. Two of them would leave that new hire lost. Fix those two, in `build_tools()`.
3. Run a ticket with the trace on and read what each tool RETURNED, not just which tools were called. An error inside a tool result is a briefing problem, not a backend problem.

```
$ python3 run.py K7PQ2M --trace
```

4. Then take the gate.

```
$ python3 verify.py 1.3
```

YOU WILL KNOW IT WORKED WHEN

Nine tools still there and none renamed, every description a real sentence, and no tool answering with an error about its own arguments. The character counts in the quote are the reference build's descriptions, not a target: yours are your own words, so expect different numbers above the floor.

```
$ python3 verify.py 1.3
```

Step 1.3: Build 1 • Make the tools route

- ✓ all 9 tools present, none renamed
- ✓ lookup_booking: description is substantive (283 chars, floor is 40), routing is proven at gate 2.1
- ✓ get_flight_status: description is substantive (218 chars, floor is 40), routing is proven at gate 2.1
- ✓ search_alternatives: description is substantive (330 chars, floor is 40), routing is proven at gate 2.1
- ✓ check_policy: description is substantive (445 chars, floor is 40), routing is proven at gate 2.1
- ✓ hold_seat: description is substantive (73 chars, floor is 40), routing is proven at gate 2.1
- ✓ confirm_rebooking: description is substantive (216 chars, floor is 40), routing is proven at gate 2.1
- ✓ issue_voucher: description is substantive (227 chars, floor is 40), routing is proven at gate 2.1
- ✓ escalate_to_human: description is substantive (220 chars, floor is 40), routing is proven at gate 2.1
- ✓ send_confirmation: description is substantive (71 chars, floor is 40), routing is proven at gate 2.1
- ✓ get_flight_status answered with data, not a refusal (0 of 1 call(s) refused)

All checks passed. Evidence code: D42-B53

That code is your receipt for step 1.3. Nothing to upload.

IF NOT YOU SEE

WHERE TO LOOK

A tool comes back with an error about one of its arguments

What did that field's description tell Claude to pass?
Read the field description, then read what the trace
shows Claude actually sent.

A tool is never picked at all

Does the description say WHEN to reach for it, or only
what the function is called? Read yours back as a
stranger.

IN CLAUDE CODE

Before you rewrite a description, not after.

```
I am about to rewrite a tool description in build_tools(). Before anything  
else, ask me what  
I think it has to say and what has to already be true before the tool is worth  
calling.  
Challenge the gaps in my answer. Do not write it for me until I have said my  
plan out loud,  
and when I am stuck point at my trace instead of at the fix.
```

NEXT

Step 1.4. Five shapes, not one.

1.4 All five shapes, numbers written down

WHAT YOU ARE BUILDING

One clean ticket passing is not the same claim as it works.

Two numbers off the totals footer are the baseline the next build has to move.
They cannot be rebuilt later.

D0

1. Run all five Stage 1 shapes with the trace on.

```
$ python3 run.py --all --trace
```

2. Read anything that failed on its own PNR, fix the one thing that trace shows, then run all five again.
3. Two of the five are not what they look like. The group of twelve escalating is a pass, not a miss. The abusive ticket comes back calm and helpful today, because there is nothing in here watching tone yet: that is a later build, so do not fix it here.
4. Read the totals footer at the end of the run and write down two numbers. Total turns, and total tokens in. In your own notes, now. The footer below is one reference build, so your turns should land close and your token total will not match it to the digit.

pnr	shape	turns	tools	in	out
stop_reason					
K7PQ2M	Clean cancellation	4	3	13,144	628
end_turn					
M3XR8T	Delay under threshold	4	3	13,043	643
end_turn					
T9WN4C	Ambiguous missed connection	4	3	13,033	674
end_turn					
G2HL9V	Out-of-scope group	3	2	9,439	701
end_turn					
R8KD3F	Abusive message	4	3	13,047	701
end_turn					

TOTAL	5 shapes	19	14	61,706	3,347
5/5 returned text					

5. Then take the gate.

```
$ python3 verify.py 1.4
```

6. At the end of the build, one person publishes the canon. This time it is the repo's creator.

```
$ python3 readout.py
$ python3 pod_sync.py --push-canon
```

7. Everyone else takes it.

```
$ python3 pod_sync.py --take-canon
```

YOU WILL KNOW IT WORKED WHEN

All five shapes generalized. Resolved is the gate's narrow word: the loop terminated and called at least one tool. It says nothing about whether the answer was any good.

```
$ python3 verify.py 1.4
```

Step 1.4: Build 1 · All five shapes

- ✓ K7PQ2M (Clean cancellation) resolved with at least one tool call
- ✓ M3XR8T (Delay under threshold) resolved with at least one tool call
- ✓ T9WN4C (Ambiguous missed connection) resolved with at least one tool call
- ✓ G2HL9V (Out-of-scope group) resolved with at least one tool call
- ✓ R8KD3F (Abusive message) resolved with at least one tool call
- ✓ 5/5 Stage 1 shapes generalized
- ✓ (R8KD3F, the abusive-message ticket, comes back calm and helpful with no gate on tone at all. That is Build 4's work, not a bug to fix here. Fixing it now hides the gap this run exists to show you.)

All checks passed. Evidence code: 46C-403

That code is your receipt for step 1.4. Nothing to upload.

IF NOT

YOU SEE

WHERE TO LOOK

One fails and four pass

That is information, not a broken loop. Read that one with the trace on its own PNR and ask what is different about that customer's shape.

The group booking reads as a failure

Read what the check actually asks. Escalation is the correct outcome for a group of twelve, so the question is whether it escalated, not whether it resolved.

IF YOU ARE AHEAD · after the gate, never instead of it

Where the context went. Open the readout you just pushed, find the turn with the largest tokens in, and write one sentence: why was that the most expensive turn, and what would have to change for the last turn to cost less than the one before it? Every turn resends the whole conversation, so the last turn is always carrying the first one.

One sentence in `PITCH.md`, in your own words, under a heading of its own.

NEXT

Build 2. Nine tools were handed to you. From here nobody tells you what belongs in the file.

Build 2 · Your own tools

Build 1 handed you nine tools already written. This one does not. You pick the tool, you write the description, you wire the function, and you defend why that one and not the other two. No client hands you a tool list on day one either.

YOU WALK OUT WITH a tool you specified yourself, a trace where Claude reached for it unprompted, that same tool served over MCP, and what each of the two cost you in context.

The gap

Nine tools, and not one of them answers what a stranded customer asks first. When is the first day I can fly?

Three questions Larkspur customers ask every day, and nothing in the file to call for any of them.

Two seams, both in agent.py

SEAM	WHAT BELONGS IN IT
EXTRA_TOOLS	The schema. Name, input shape, and the description string Claude chooses from.
LOCAL_TOOLS	The function, registered under exactly that name. A schema with no function is a tool the loop cannot dispatch.

Three candidates

Every candidate is grounded in a document the client already owns. A tool whose answer you cannot point at invents flight facts, and that is gate one.

Pick one. Two or three is the ceiling, not the target.

CANDIDATE	WHERE ITS ANSWER COMES FROM
next_available_day	Live inventory in <code>data/americas/</code> . The wrapper function is given, and the <code>LOCAL_TOOLS</code> line for it is there, commented out. <code>EXTRA_TOOLS</code> is an empty list, so the whole schema is yours to write: name, input shape and description.
fare_rules	The Handbook excerpt in the repo: fare family, change fee, fare difference.

2.1 Your own tool, chosen unprompted

WHAT YOU ARE BUILDING

A tool the model never picks is a tool you never shipped.

Every schema you register is resent on every turn, fired or not. One new capability is a permanent tax, and you are about to price it.

DO

1. Argue one tool off the shortlist before anyone types. Five minutes of talking buys you thirty of building.

```
$ python3 run.py --show-tools
```

2. Write the schema into `EXTRA_TOOLS`, then wire the function and register it in `LOCAL_TOOLS` under exactly that name. Declared and undispatchable is the most common way this build breaks.
3. Check the wiring before you spend a probe on it. Yours is listed beside the given nine, tagged `+ yours`, with the description you actually wrote.

```
$ python3 run.py --show-tools
```

4. Write your pod's probe. Three lines in `build2_probe.txt` at the repo root: PNR, last name, the customer's message. One question for the whole pod, so you test one thing instead of five.

```
K7PQ2M  
Marchetti  
When is the first day I can actually fly?
```

5. Write the probe before the description is finished. One reverse-engineered from a tool that already fires proves nothing.
6. Run the five shapes again and put tokens in beside the two numbers you wrote down in Build 1. Up or down, both are real answers.

```
$ python3 run.py --all --trace
```

7. Then take the gate.

```
$ python3 verify.py 2.1
```

YOU WILL KNOW IT WORKED WHEN

The gate does not care that the tool exists. It cares that the model chose it, on a message that never named it. Three attempts, one success is enough, and declared-but-never-called does not pass.

```
$ python3 verify.py 2.1
```

Step 2.1: Build 2 • Your own tool

- ✓ at least one tool beyond the given nine (next_available_day)
- ✓ the new tool is registered in LOCAL_TOOLS
- ✓ next_available_day description is $\geq$ 40 characters (628)
- ✓ the model chose a new tool on a conversation that needs it (next_available_day, attempt 1 of 3)
- ✓ returned a non-empty string
- ✓ (that probe cost 1,205 tokens in across 2 tool call(s), the tenth tool is not free, and this is the number that says what it cost)

All checks passed. Evidence code: 830-C95

That code is your receipt for step 2.1. Nothing to upload.

IF NOT

YOU SEE

It never gets picked and you cannot see why

It only fires when you name it in the message

WHERE TO LOOK

If `--show-tools` lists yours tagged `+ yours`, the wiring is in and the suspect is the description. Go there first, and go there before you touch the probe.

Read your probe as the model sees it. Does that customer's message need your tool, or could one of the given nine answer it?

IN CLAUDE CODE

Once your pod has argued one tool off the shortlist, before you type the schema.

```
I'm adding a tool to this agent. Draft its description with me line by line, asking me why each clause is there and where its boundary sits against search_alternatives. Do not write the description for me, and do not touch EXTRA_TOOLS until I've said the wording myself.
```

IF YOU ARE AHEAD · after the gate, never instead of it

The second tool. Add `fare_rules` or `reopen_stats`, run all five shapes again, and answer two questions off your own trace. With the given nine plus every one of yours going up on every call, did routing still hold on all five shapes, and what did tokens in per contact do? Past about fifteen tools routing starts to slip and the token tax climbs. End it with keep-it-or-delete-it.

The keep-or-delete line in `PITCH.md`, with the number that decided it.

NEXT

Step 2.2. Same tool, same schema, off your laptop.

2.2 The same tool, over MCP

WHAT YOU ARE BUILDING

Same schema, different transport. The tool stops living in your file and starts living in a program of its own, versioned on its own, reachable by any host.

The client asks the server what it has, Claude picks, the client calls it.

`tools/list`, then `tools/call`.

DO

1. Run the selftest and watch the two messages that go past. `tools/list` is the server handing over every tool it has, with its JSON schema. `tools/call` is one of them being run.

```
$ python3 support/mcp_selftest.py
```

2. In `agent.py`, at the seam marked `2.2`, read the question there and answer it in code. The dispatch is given: a call for a name the server owns already goes to the server. What is not given is what the model is offered, and what has to leave the file for the server to be the only owner.
3. Check the wiring. Yours should now be tagged `via mcp`.

```
$ python3 run.py --show-tools
```

4. Run your probe's question with the trace on and find `[mcp]` on the call. Same tool, same arguments, different owner.

```
$ python3 run.py K7PQ2M --trace --message "When is the first day I can actually fly?"
```

5. Put tokens in beside your 2.1 number. It moves by what the server offers, not by the transport. On the reference run it went up 989, about 10%, and the extra is the second tool the server happens to carry coming along with the merge. That is the teach: every schema you are offered rides on every turn, wherever it lives. MCP changes what you maintain, and does nothing at all for routing or for what the tool list costs you.
6. Then take the gate.

```
$ python3 verify.py 2.2
```

YOU WILL KNOW IT WORKED WHEN

A tool discovered over MCP fired on your pod's probe, and the local copy is gone. One tool, one owner. The last line is the gate showing its work on the token question: it names both numbers and the band it read them against.

```
$ python3 verify.py 2.2
```

Step 2.2: Build 2 · The same tool, over MCP

- ✓ the MCP server answers tools/list (fare_rules, next_available_day)
- ✓ at least one MCP tool is offered to Claude (fare_rules, next_available_day)
- ✓ no tool name is served twice, once here and once over MCP (clean)
- ✓ an MCP-discovered tool fired on the pod's probe (next_available_day, attempt 1 of 3)
- ✓ returned a non-empty string
- ✓ (tokens in: 10,320 at 2.1, 11,309 here, +989: inside the +/- 1032 noise band. Same schema, same tax. A jump up means the tool list GREW rather than changed hands; --show-tools says which, in one screen.)

All checks passed. Evidence code: 662-1D7

That code is your receipt for step 2.2. Nothing to upload.

IF NOT

YOU SEE

WHERE TO LOOK

The tool shows up twice

The local copy is still registered. One name, one owner.

Nothing is tagged via mcp

Is the server process starting at all? Everything it wants to say to a human goes to stderr, so read stderr.

It fires, but tokens in went up

Count the schemas on the wire. How many tools went up on that turn, and is one of them the same tool twice?

IF YOU ARE AHEAD · after the gate, never instead of it

One server, many hosts. Register the same server in Claude Code, ask it when J5NU8S can fly, and watch a host you did not write drive your server. Then take it out again.

```
claude mcp add larkspur -- python3 support/mcp_server.py, then claude mcp remove larkspur when you are done.
```

NEXT

Pitch and ship. The build stops and the claim starts.

Pitch and ship

The deliverable is a repo, not a recording. Before it goes anywhere, six lines that say what you built and what it is worth. One of those six names what is broken.

The pitch, six lines

"Compared to what?" is the question that separates a demo from a claim. Which of your six survives it?

LINE	WHAT IT CARRIES
Built:	What we built, and what you connected.
Does:	What it does for a stranded customer, in one line.
Number:	Before and after, with a unit, off your own trace, and what you measured it on.
Guardrail:	The one you can prove, and how you proved it.
Next:	What you would do with another week.
Still broken:	The gap you name before anyone finds it. Scored like the other five. The gate reads this label by name.

Three things are provable at this point

Accuracy and dollars are not on that list. There is no eval suite yet, so any number about quality or cost per contact is one you cannot defend. Saying so out loud is stronger than guessing, and it is exactly what the next session builds.

- Your Stage 1 pass rate.
- Your before and after turn count and context cost, measured by you.
- A forged confirmation token being refused, while a real customer click is accepted.

The template, then make every line yours

This is the shipped `PITCH.md`, byte for byte. Fill every line after the colon. `Lever:` stays as the menu until Build 4 picks one word; the gates read `Still broken:` and `Lever:` by name, so keep both labels exactly.

What each line wants. Built: what we built, one line. Does: what it does for a stranded customer. Number: before to after, with a unit, measured on what. Guardrail: the one we can prove, and how. Next: what we would do with another week. Still broken: the gap we are naming before anyone finds it. One typist per pod: merge conflicts eat the whole thing.

```
# PITCH.md
```

```
Six lines and a lever. Both of the last two are scored.
```

```
Built:
```

```
Does:
```

```
Number:
```

```
Guardrail:
```

```
Next:
```

```
Still broken:
```

```
Lever: <cost | speed | intelligence>
```

Publish the pod's canon

This build's committer goes first, and tags last, once the canon is on the remote. The readout is generated from the agent being published, so it describes the version the pod is about to stand on.

```
$ python3 readout.py
$ python3 pod_sync.py --push-canon
$ git tag v1 && git push origin HEAD --tags
```

Then everyone else takes it, once the committer has said it is up. Take a canon nobody has pushed yet and you pull down the last build's agent.

```
$ python3 pod_sync.py --take-canon
```

Prove the key never travelled

The first prints `.env`. The second must print nothing at all. Anything it prints is an incident rather than a tidy-up: stop, and say so out loud.

```
$ git check-ignore .env && git log --all --oneline -- .env
```

And this is what gets run on it

The rubric is public and it runs on the shared screen. Overclaiming scores worse than a gap.

Nothing is required between the sessions. The next session starts from the pod's canon, not from your laptop's memory of it.

DIMENSION	WHAT EARNS THE 2
Runs	The loop gate and the generalization gate both banked.

Generalizes	All five Stage 1 shapes, not the easy one.
Tools	Every description says when to call it and what comes back.
Guardrail	Tested, not judged. A forged token is refused.
Claim	A number you measured, and a claim sized to it.

Between sessions

Nothing is required of you. One file lands in your repo, and the next session opens on it, so read it when it arrives.

A bot read your build

It read your code, not your demo. It ran nothing, and it changed exactly one file: `FEEDBACK.md`, at the root of your own repo.

It is written into the pod repo before the next session opens, so it is not in a fresh clone and it is not something you generate. If `git pull` does not bring it down, ask in your pod thread rather than hunting for it.

```
$ git pull
$ open FEEDBACK.md
```

What is in it

SECTION	WHAT IS IN IT
Priya's note	Her push, in her words: use your best model.
Still open	The same four questions, with what she means by each.
What the review agent found	Three to five observations, in your own file and function names.
Before our 9:00	Which model should we be on, and how will you prove it.

Her four questions

None of the four is answered by naming a model. A claim with no number under it is an opinion, and her vendor already supplied one.

SHE ASKED	WHAT SHE MEANS BY IT
What it costs	Per resolved contact, against the \$6.90 a human contact costs them.
When it is wrong	The first untrue thing it says, and what happens after that.
Who runs it	In June, after you have left.
What you left out	The scope you cut, and why.

First thing, next session

Before anything else, on every laptop. What your pod shipped is on the remote, tagged `v1`. What is on your laptop is a stale guess. Take it, run it once, and the pod is arguing about one agent again.

It saves your own `agent.py` into `.workshop/mine/` first and prints where it went. Nothing you wrote is lost.

```
$ python3 pod_sync.py --take-canon  
$ python3 run.py K7PQ2M --trace
```

Build 3 · Build the proof

A claim is only as good as its instrument. You made a claim in the pitch and nothing has graded it. This build makes the thing that can, and the cases carry your name because that is the answer to who else can operate this.

YOU WALK OUT WITH an eval suite with a case authored by you, a surface a sponsor can watch, and a result you would read out loud even when it says **BLOCKED**.

The four pieces

PIECE	WHAT IT IS
The case	An input and what you expect back. Frozen, so it means the same thing later.
The run	Your agent, on that input, unmodified.
The grader	The thing that decides pass or fail, and the half everyone forgets is code too.
The gate	Whether a failure blocks a release on its own. Hard, or soft.

Which grader to reach for

Two habits before you ship a judge. Make it quote what it read, and never let a model grade its own output.

And suspect the eval first. A failing eval is a claim about two things, and only one of them is your agent. Larkspur's week 8 lost a suite to 37 of 48 in a single night, the only change was the rubric, and nine of eleven failures were the grader. Fixed in 40 minutes, agent untouched. It happened on this exercise's own first live run too: read `evals/GRADER-BUG.md` in your repo.

GRADER	GOOD FOR	CHARGES YOU
Exact match	A PNR, a flight number, a policy row id	Nothing, and it fails a right answer worded differently
Rule-based	Was <code>confirm_rebooking</code> called? Was a hotel promised? Roughly seven in ten production checks live here	You have to say the rule out loud

LLM as
judge

Tone, acknowledgement, whether an explanation is
honest

Money, latency, and a
grader that can itself be
wrong

3.1 Build the proof

WHAT YOU ARE BUILDING

Three surfaces, three audiences. The trace is for you, the chat is for the customer, the panel is for the sponsor. Most teams build the first, demo the second, and never build the third.

A blocked release can pass this one. Honest proof beats a green dashboard.

DO

1. Open the surface. There is nothing to wire: `python3 demo/serve.py` serves it as shipped, on `http://localhost:4390`, with chat on the left, the evidence panel on the right, and your own `run_agent` behind both. It is given on purpose: the proof is the work, not the CSS.

```
$ python3 demo/serve.py
```

2. Copy the worked examples, then make them yours. A pod that ships them untouched does not clear the gate.

```
$ cp evals/cases.example.json evals/cases.json
```

3. Write one case per member, three minimum, each carrying its author's name in the `author` field. That field is matched against the name banking the gate, so a pod of six cannot ship six cases with one person's judgment in all of them.

4. Make two of them hard gates. A hard gate is a case whose failure blocks a release on its own. If everything is soft, nothing is gated.

5. Run the suite against your own agent.

```
$ python3 eval_harness.py
```

6. Fill the panel. Pass rate, what blocked, one claim, the caveat. Behind on time, cut the panel: the cases are the proof, the panel only puts them on a screen.

7. Then take the gate. It does not require your cases to pass.

```
$ python3 verify.py 3.1
```

8. Bring three things to the share-out. The case you wrote first and the customer shape behind it, how you graded it and the sentence your expectation turns on, and the one your agent failed. Say that third one out loud: it is the most useful minute of the build.

9. At the end of the build, the third name on the roster publishes the canon. Everyone else takes it.

```
$ python3 readout.py  
$ python3 pod_sync.py --push-canon
```

YOU WILL KNOW IT WORKED WHEN

Ten checks, and every label names the file it read. Seven read your cases and your last run: at least 3 cases, at least 2 hard gates, every case says what it expects, one authored by you, the last run was YOUR cases, the run covered at least 3 cases, the panel has numbers. Three read `PITCH.md`: a line carries a number, the file is more than the shipped template, and the `Still broken:` line names one thing. Whether your cases passed is not on that list, and never will be. A blocked release passes this gate. Bracketed numbers and the evidence code are one reference run's; yours will differ.

```
$ python3 verify.py 3.1
```

Step 3.1: Build 3 · Build the proof

- ✓ at least 3 eval cases (5)
- ✓ at least 2 of them are hard gates (2)
- ✓ every case says what it expects
- ✓ at least one case is authored by you (Alice Example)
- ✓ the last `eval_harness` run was YOUR cases, not the examples
- ✓ the run covered at least 3 cases (ran 5 of the 5 in `cases.json`, 5 scored, release BLOCKED)
 - ✓ the evidence panel has numbers to show
 - ✓ `PITCH.md`, whole file: some line carries a number (the `'Number:'` line is where it belongs)
 - ✓ `PITCH.md`, whole file: more than the shipped template (138 words, floor is 40)
 - ✓ `PITCH.md`: the `'Still broken:'` line names one thing that still does not work

All checks passed. Evidence code: 8E5-8BB

That code is your receipt for step 3.1. Nothing to upload.

IF NOT YOU SEE

WHERE TO LOOK

Every case passes

You wrote them too gently. Ask what you would write if you were trying to find the thing that breaks, then write that one.

A case fails on behaviour you believe was right

Suspect the rubric before the agent. Read `evals/GRADER-BUG.md`, which is that exact failure from Larkspur's own week 8.

A row reads `UNKN` and the run says the grader could not read its own judge

That case ran and was not scored, and it is not charged to your agent. The pass rate is over scored cases only. Run the suite again; if it repeats, the judge prompt for that case is the suspect, not the agent.

IN CLAUDE CODE

Once you know which customer shape your case is about, before you fight the JSON.

```
Write the eval case I'm about to describe, then argue against it as the grader would:
every way my expect could pass a bad answer or fail a good one. Ask me which grader I want
before you suggest one, and tell me what that grader charges me.
```

IF YOU ARE AHEAD · after the gate, never instead of it

Agent, or grader? Take one failing case and decide, in writing, which one it was. If it was the grader, change only `expect`, bump `rubric_version` to `v3`, keep the v2 file beside it, and run it again. A verdict that flips with `agent.py` untouched is your own week 8.

Two rubric files, v2 and v3, side by side, and the verdict that flipped between them.

NEXT

Build 4. Now something in your repo can read a saving, you can go and make one.

Build 4 · Pull the lever

Proof came first for a reason. You cannot bank a saving you have no way to read. Now you move one number on purpose and prove it moved, without breaking anything that already worked.

YOU WALK OUT WITH one declared lever, a before and after pair off your own laptop, and a movement you would put in front of somebody who owns the budget.

What each lever buys, and what it charges

"All three" is a dodge. Which one first, and why that one?

LEVER	BUYS	CHARGES
Intelligence	Harder shapes, fewer wrong answers	Tokens, latency, more to evaluate
Cost	A number the sponsor can defend	A cache you must keep valid, and latency: p50 and p95 both went up about a fifth on the reference run
Speed	Abandonment, and the SLA	A smaller model. A quality bet.

Your lane, and where to look

The intelligence lane is allowed to be a different model, one tier up or one tier down, as long as you bench the same stage on both sides of the swap.

LANE	YOUR NUMBER	WHERE TO LOOK
cost	Model \$ per contact	Cache, tokens, the tool wall
speed	p95	Triage, parallel reads, turns
intelligence	Wire rules on Stage 2	Five adversarial messages, one tier either way

Real measurements, and not one of them a model upgrade

\$0.14 per resolved contact against \$6.90 for a human chat is the only comparison a board reads, and it flatters. Humans still take 42% of in-scope chats, and all voice. Say the caveat, and the date the prices came from, on the same slide.

Caching can also cost you more than it saves. Writing to the cache is priced above fresh input, so a prefix that changes on every call pays the premium and never collects. Larkspur read that as caching not working here, and carried the premium for two weeks.

`bench.py` reports model cost only. Larkspur's loaded cost per resolved contact was about \$0.14 against a model cost of \$0.087, so roughly 40% of the real number is infrastructure and evals. Report it as model cost, or say loaded and estimated. Do not quietly call it the loaded number.

CONFIGURATION	WEEK	P95	MODEL \$/CONTACT
One mid-tier model, no cache	2	9.8s	\$0.173
plus caching, parallel reads	5	6.9s	\$0.102
plus small-model triage, trimmed output	7	4.7s	\$0.087

4.1 Pull the lever, measure it

WHAT YOU ARE BUILDING

Bench before you tune. The before number cannot be rebuilt after.
A win that breaks a Stage 1 shape fails, and that is the point of the build.

DO

1. Declare your lane in writing, in `PITCH.md` and in your pod thread. The shipped file carries the line as a menu, `Lever: <cost | speed | intelligence>`. Replace everything after the colon with the one word you picked, because the gate reads that word and refuses the menu.

```
Lever: cost
```

2. Every member benches cold on their own laptop, labelled `before`, before any change lands. A podmate's fix does not carry your before number, and it cannot be rebuilt. Three runs per shape, both sides: at one run per shape an output-token delta under about 15% is noise, and `--compare` prints it beside your win as if it were real.

```
$ python3 bench.py --label before --stage 1 --runs 3
```

3. Intelligence lane only: bench Stage 2 on both sides of the swap as well. Your metric is the wire-rule count, and that only exists on Stage 2.

```
$ python3 bench.py --label s2-before --stage 2 --runs 3
# then, after the swap
$ python3 bench.py --label s2-after --stage 2 --runs 3
```

4. Pull one lever. The one you declared.
5. Bench again and diff the two labels. Same stage and same number of runs on both sides, or the comparison is not one.

```
$ python3 bench.py --label after --stage 1 --runs 3
$ python3 bench.py --compare before after
```

6. Walk the five shapes again before you call it done. A win you did not recheck is a claim you cannot make.

```
$ python3 run.py --all --trace
```

7. Then take the gate.

```
$ python3 verify.py 4.1
```

8. At the end of the build, the fourth name on the roster publishes the canon. The last hand. Everyone else takes it.

```
$ python3 readout.py
$ python3 pod_sync.py --push-canon
```

YOU WILL KNOW IT WORKED WHEN

Your declared lane's number moved at least 10% in the right direction, and Stage 1 still resolves 5/5. Ten percent is there so noise does not read as a win. A 70% cost win that breaks a shape fails. The first check reads the `Lever:` line in `PITCH.md`. The numbers below are one reference run on the cost lane; yours will differ.

```
$ python3 verify.py 4.1
```

Step 4.1: Build 4 • Pull the lever

- ✓ `PITCH.md`: the `'Lever:'` line names one lane (cost)
- ✓ a `'before'` bench exists
- ✓ an `'after'` bench exists
- ✓ the two runs are comparable
- ✓ stage 1 still resolves 5/5 after your change (100%)
- ✓ cost lane: `model_cost_per_contact` 0.0559 → 0.0234 (-58%)

All checks passed. Evidence code: 61A-096

That code is your receipt for step 4.1. Nothing to upload.

IF NOT

YOU SEE

WHERE TO LOOK

The number went the wrong way

Read the bench file before you touch the agent again. Cache read and written, tokens, turns. The row that moved is not always the row you aimed at.

The bench says `cache: not used on any turn`

Nothing has been marked for reuse, so a stable prefix has nothing to hit. That is not a 0% hit rate; it is no cache at all. What would the API have to be told, about the part of the request that does not change, before it can keep anything between calls?

Caching reports a 0% hit rate

Something is being kept and nothing is ever found again. Print what goes out at the very top of every request, two calls in a row, and ask which characters are not identical.

Cost fell and latency rose

That is the cost lever's real charge: the reference run paid about a fifth on p50 and p95. Say it in the pitch. A lever that only shows its buys is a claim the client will find the charge on.

IN CLAUDE CODE

After the before bench, while you are deciding which lever to pull.

```
Read .workshop/bench-before.json and my latest trace, and tell me where the
tokens go per
turn. Ask me which lever I think that points to, and why, before you say
yours.
Do not change agent.py.
```

IF YOU ARE AHEAD · after the gate, never instead of it

The second lever. Leave `before` and `after` untouched, they are your banked pair. Pull a second lever, bench it as `after2`, and compare `after` against `after2`. The payoff is not the row that improved. It is the row that got worse, with its percentage. That is the Charges column, on your own build.

The row that got worse, named with its number, in `PITCH.md`.

NEXT

Pitch the client. Somebody who owns a budget is about to ask you what a year of this costs.

Pitch the client

No neutral audience. Whoever is not pitching is the client, and clients own budgets. You are paired across lanes, so the client opposite you moved a different number and knows exactly where your claim gets thin.

How the twenty five minutes run

PHASE	MIN	WHAT HAPPENS
You pitch	5	Demo the surface, state the claim, name the caveat
They are the client	4	Two questions, off the four axes. Not four soft ones.
Swap, they pitch	5	Same three beats
You are the client	4	Two questions back
Everyone writes	3	Your three lines, filed as an issue on the repo you just heard
Regroup	4	Back in the main room once your issue is filed

Three things go in with you

Claim only what your evidence shows. "We did not measure that" beats a guess, and the caveat is scored, so a pod that hides a gap loses more than a pod that owns one.

- One claim. A number, its unit and its denominator, on the one metric your pod declared.
- Two numbers. Before and after, straight off your own two bench files.
- One caveat. What you did not measure, named by you before she finds it.

Pick a persona, ask two questions

Holding a fresh cost or speed number? Draw Marcus. He is the harder room, and that is the point. His bar, in his words: what does a year of this cost, and whose name is on it?

Your questions are graded too. Four soft ones and the room learns nothing. Ask them the question you would least like asked of your own build.

THE AXIS	PRIYA • PROJECT MANAGER	MARCUS • HEAD OF OPS
cost	Per resolved contact	Per year, at our volume
wrong	The first untrue thing it says	Who is accountable when it is

operate

Who runs it in June

The day your team leaves

scope

What you left out, and why

Where the kill switch lives

Three lines, filed where the work lives

A number you would take to your own client, a claim you would not sign, and a gap. Filed as a signed issue on their repo, where the build lives, rather than in a chat that closes when the session does.

Then read the one on yours before the close. Every pod leaves with a read on their work that nobody running the room wrote.

What is provable now

Be precise about the claim. That precision is the last thing this build teaches, and it is the half that travels.

Provable, and not yet

That is the honest state of most week-four engagements.

PROVABLE	NOT YET
Three gated eval cases, authored and run	Accuracy at volume
Your own before and after on a metric you chose	Real customer traffic
A forged confirmation token refused	Anything at production scale

Two sessions, one agent

SESSION ONE · DONE	SESSION TWO · DONE	FROM HERE · YOURS
Build. The harness, the loop, and tools you wrote yourself.	Prove. Eval cases, a grader, and a number that moved.	Scale. Nobody in this room does that part for you.

The proof pack, before you go

The named account is the ask. A repo nobody takes to a customer inside 30 days is a training exercise. Not a segment, not a vertical: a logo, a workflow you understand end to end, and a date on a calendar.

You keep the repo, the eval cases and the grader. The grader is the half that travels, and it is the cheapest rehearsal available before a real pitch.

WHAT YOU RECORD	WHERE IT COMES FROM
The repo URL	Tagged <code>v1</code> when you shipped it.
The six <code>PITCH.md</code> lines	Written at the pitch, sharpened when you pitched the client.
Your two bench numbers	<code>bench.py</code> , before and after.
One named account, its use case, and a date inside 30 days	You. Nobody else can write this line.

Anything you took past a gate

The stretch box on your build steps. Not scored.

A · Working in one repo

The rule, in one sentence. **Nobody commits** `agent.py` **mid-build. At the end, the build's committer publishes the canon and everyone else takes it.**

```
$ python3 pod_sync.py --push-canon # the committer
$ python3 pod_sync.py --take-canon # everyone else
```

Everyone builds their own `agent.py` on their own laptop. Several people editing one file in one repo at the same time produces a merge conflict inside a loop they are all still learning to read, and no breakout has time for that. So the repo only ever holds one version: the one the pod chose.

`--take-canon` saves your version into `.workshop/mine/` first and prints the path. Nothing you wrote is lost, and diffing yours against the canon is the most useful five minutes of the build.

`PITCH.md` and `evals/cases.json` are the pod's shared record and they are normal commits. The bench pair never leaves your laptop.

Four hands, in order

BUILD	WHO PUBLISHES THE CANON
Build 1	The repo's creator
Build 2	The second name on the roster
Build 3	The third name
Build 4	The fourth name, the last hand

Behind, or something broke

The pod repo is the checkpoint. This restores what the pod pushed.

```
$ python3 pod_sync.py --take-canon --force
```

What it does not restore

If `.workshop/` is already gone, say so rather than re-benching `before` against the current agent. That number would be a fiction, and your pod's claim rests on it.

NOT RESTORED	WHAT TO DO

`.env`, your key

Yours, and never in the repo. Re-create it from

`.env.example`.

`.venv/`

Rebuild it with `python3 setup.py --fix`.

`.workshop/`, your banked gates
and bench numbers

Not on the remote. Copy it across before you re-clone into
a fresh folder, and copy it back in after.

B · The five tickets

One example of the bar, not the only wording that passes. Two of the five are not what they look like: one should not resolve at all, and one gets an answer this build has no gate for.

PNR	SHAPE	WHAT GOOD LOOKS LIKE
K7PQ2M	Clean cancellation	Holds a seat, then waits for the customer's click
M3XR8T	Delay under threshold	Refuses the waiver and names the policy row it came from
T9WN4C	Ambiguous missed connection	Asks a clarifying question before it does anything
G2HL9V	Out of scope, a group of 12	Escalates. It does not resolve, and that is the pass
R8KD3F	Abuse plus a legal threat	Comes back calm and helpful. There is no tone gate yet

C · The words of the build

WORD	WHAT IT MEANS HERE
PNR	The booking reference, six characters. It is how every tool finds this customer.
Turn	One request plus one reply. You resend the transcript, and that IS the memory.
<code>stop_reason</code>	The field that ends a turn or continues it. <code>tool_use</code> means Claude is waiting on an answer, <code>end_turn</code> means it is done.
<code>tool_result</code>	The block you send back with what a tool returned, keyed to the id of the request that asked for it.
Tool description	The only thing Claude reads to pick a tool. A briefing, not docs.
Hard gate	One failure blocks the release. Never an average.
Canon	The one <code>agent.py</code> the pod agreed on, published at the end of a build. Until you take it, your laptop is a guess.
Probe	Your pod's one test message for a tool you wrote, three lines in <code>build2_probe.txt</code> : PNR, last name, the customer's message.
Bench	One timed run of the agent over a fixed set of tickets, labelled so two runs can be diffed.
p95	The response time 95 of every 100 contacts come in under. It is the number an SLA is written against, not the average.

D · Larkspur's systems as the agent sees them

Every line here is sourced from `support/tools.py` and `data/americas/disruption_policy.json`. This is the map of what is fenced, so nothing is written into it that the code does not do.

Reads

TOOL	WHAT IT DOES
<code>lookup_booking</code>	Takes a PNR and a last name, and refuses if the name is not on that booking. Returns a trimmed view, never the raw record, and labels the ops note and the remarks as untrusted free text.
<code>get_flight_status</code>	Takes a flight number and a date. The one read not fenced to the customer's own booking, so it is the one that can be pointed anywhere.
<code>search_alternatives</code>	Takes the PNR alone, so it can never be pointed at a route the customer did not buy.
<code>check_policy</code>	Re-derives fare family, loyalty tier and whether this is an overnight from the booking on every call, so a model cannot talk its way into an entitlement the data does not support.

Writes

TOOL	WHAT IT DOES	THE FENCE
<code>hold_seat</code>	Puts a seat aside and hands back a hold that expires in 15 minutes.	Reversible, and it undoes itself if nobody confirms.
<code>issue_voucher</code>	Issues on its own only under the policy threshold: 25 dollars for a meal, 40 for ground, 75 for goodwill.	A hotel is never automatic, and anything over the line queues for a human.
<code>confirm_rebooking</code>	Reissues the ticket. The irreversible one.	Only with the customer's own click token. "The customer said yes" in chat is not it.
<code>send_confirmation</code>	Writes the message the customer actually reads.	The only write with no threshold and no token, so it carries whatever the agent got wrong.

escalate_to_human

Files the case with a written summary for whoever picks it up.

The correct outcome for groups, refunds, and anything the policy does not cover.

BASECAMP

Prepared for Partner Basecamp participants. Not for reproduction or redistribution as training material. You are free to apply these patterns in your own client work. Larkspur Airlines is fictional; data/americas/ is the authority for every fact the agent states.